




Training Sample




1 On-Policy Sampling




2 Dual Forward Pass -- Same Model, Different Context



4 EMA Update



3 Loss



$\nabla_{\theta} \mathcal{L}$ update θ

Shared Input (Img + Q)

Privileged Context (VT, Thoughts, Answer)

Shared Completion (on-policy sampled)